

Concurrency in C++11 series

pabloxrl@gmail.com

Introduction

- Purposes:
 - Bring everyone up to speed with C++11 Multithreading/Concurrency facilities
 - Provide insightful examples on how to use new C++11 features
- Different topics will be covered in this series incrementally:
 - Basic thread management
 - Thread-safe data structures
 - Lock-free data structures
 - And more

Introduction

- Baseline (or what do I assume you know):
 - What is concurrency
 - How multithreading differs from task switching
 - The difference between concurrency with multiple processes and concurrency with multiple threads
 - Why (and why not) to use concurrency
- Sources:
 - Main source: C++ Concurrency in Action
 - Other sources will be quoted
- Code will be available in my github repository;
 - <https://github.com/pabloribalta/Concurrency.git>

C++11: Concurrency and multithreading

- The 1998 C++ standard doesn't acknowledge the existence of threads
 - Compilers had to ship extensions to enable multithreading
 - Prevalence of C APIs for multithreading
 - POSIX C standard
 - Microsoft Windows API
 - General purpose C++ libraries like Boost and ACE tried to conceal platform-specifics
- Ultimately: No support in the standard

C++11: Concurrency and multithreading

- C++11 includes a new thread-aware memory model and the C++ Standard Library has been extended
 - Boost Thread Library as a primary model (names, structure, and such)
- This comes at a cost of *abstraction penalty* but low level facilities are also provided

Hello, concurrent world

```
#include <iostream>
#include <thread>

void hello()
{
    std::cout << "Hello concurrent world\n";
}

int main()
{
    std::thread t(hello);
    t.join();

    return 0;
}
```

Hello, concurrent world (s1t01)



A diagram with four blue circular callouts containing numbers 1, 2, 3, and 4. Callout 1 points to the word 'concurrent' in the title. Callout 2 points to the opening angle bracket of the <thread> include directive. Callout 3 points to the opening curly brace of the main function. Callout 4 points to the 't(hello)' part of the thread creation line.

```
#include <iostream>
#include <thread>

void hello()
{
    std::cout << "Hello concurrent world\n";
}

int main()
{
    std::thread t(hello);
    t.join();

    return 0;
}
```

Basic thread management

- Launching threads

- Boils down to constructing an `std::thread` object

```
void do_some_work(){ ... }
```

```
std::thread my_thread(do_some_work);
```

- `std::thread` works with every *callable* type, so it can accept an instance of a class with a function call operator instead

(s1t02)

- Lambda functions can also be used, which can help to avoid the common problem dubbed as the „C++ most vexing parse”

(s1t03)

Basic thread management

- Waiting for threads to complete
 - Explicitly choose what to do with the thread
 - Wait for it to complete (`myThread.join()`)
 - Let it run free (`myThread.detach()`)
 - Failing to do so before the `std::thread` object is destroyed will result in a crash (`std::thread` destructor calls `std::terminate`)
- When not waiting for the thread completion, we must ensure that the data accessed by the thread is valid until the thread has finished with it. (sito4)

Basic thread management

- Waiting in exceptional circumstances:
 - Avoiding termination when an exception is thrown using a try/catch block (s1to5)
 - Using RAI to wait for a thread to complete (s1to6)
 - If it is not necessary to wait for a thread to finish, we can *detach* from ensuring that `std::terminate()` won't be called when `std::thread` is destroyed.

Basic thread management

- Passing arguments to a thread
 - **Common pitfall:** By default, arguments are *copied* into internal storage, where they can be accessed by the newly created thread of execution, even if the parameter should be a reference. (s1to7 and s1to8)
 - Semantics are similar to `std::bind` when using member functions. (s1to9)
 - It is possible to use `std::move` when passing non-copyable objects as parameters (s1t10)

Basic thread management

- Transferring ownership of a thread
 - Keeping in mind that `std::thread` is *movable* but not *copyable* (sit11)
 - Scenario #0: A function that creates a thread to run in the background and passes back the ownership to its caller (sit12)
 - Scenario #1: Create a thread and pass its ownership in to some function that should wait for it (sit12)
 - Possible to improve our ThreadGuard (sit13)

Basic thread management

- High level thread management
 - Determining number of runnable parallel threads with `std::thread::hardware_concurrency`
 - Identifying threads with `std::thread_id` (sit14)
 - **Bonus:** A naive parallel version of `std::accumulate` (sit15)

Basic thread management

- Summary
 - Starting threads
 - Waiting (and *not waiting*) for threads to finish
 - Running threads in the background
 - Passing arguments to threads
 - Transfer of ownership of `std::thread` objects
 - How to manage a group of threads

Next presentation

- Sharing data between threads
 - Avoiding race conditions
 - Using `std::mutex` and `std::lock`
 - Structuring code for protecting shared data
 - Flexible locking strategies
 - Transferring ownership of locks and mutexes between scopes
 - Alternative facilities for protecting shared data

The end

- Any questions?