

Concurrency in C++11/14 series

pabloxrl@gmail.com




Index

Problems with sharing data between threads

- Problems arise due to the consequences of modifying data shared between threads (read-only is ok)
- Handy concept: *Invariants*
- Example of invariant: For doubly linked list
 - If you follow a “next” pointer from one node (A) to another (B), the “previous” pointer from that node (B) points back to the first node (A).

Deleting element from doubly linked list

- Steps for deletion:
 - Identify node to delete (N)
 - Update link from N-1 to N+1
 - Update link from N+1 to N-1
 - Delete N
- Unless something is done other threads could see the list in an inconsistent state



Problems with sharing data between threads

- Race conditions are the biggest threat
 - Outcome depends on the relative ordering of execution
 - May lead to broken invariants
 - Ultimately can cause *undefined behavior*
- Ways to avoid race conditions
 - Wrap data structures with protection mechanisms
 - Modifications are done as a series of indivisible changes preserving invariants (lock free programming)
 - Handle updates as transactions

Protecting shared data with mutexes

- Using a mutex → *MUT*ually *Ex*clusive
 - 1- Lock the mutex
 - 2- Do stuff
 - 3- Unlock the mutex
- Using an instance of `std::mutex`
 - 1- Call `std::mutex::lock()`
 - 2- Do stuff
 - 3- Call `std::mutex::unlock()`

Protecting shared data with mutexes

- Using a mutex → *MUT*ually *Ex*clusive

- 1- Lock the mutex
- 2- Do stuff
- 3- Unlock the mutex

- Using an instance of `std::mutex`

- 1- Call `std::mutex::lock()` Not recommended
- 2- Do stuff
- 3- Call `std::mutex::unlock()`

Protecting shared data with mutexes

- Using mutexes in C++14
 - Standard provides `std::lock_guard<T>` class template which implements RAI (scoped locking/unlocking on supplied mutex) (s2t01)
 - `std::lock_guard` and `std::mutex` are both declared in `<mutex>` header
 - **Potential danger:** Returning references to protected data or passing functions that access protected data. (s2t02, s2t03)

Protecting shared data with mutexes

- Neat tricks for avoiding deadlocks
 - Locking more than one object at once (s2to4)
 - Flexible locking with `std::unique_lock` (s2to5)
 - Transferring ownership of locks (s2to6)
- Bonus:
 - Thread-safe lazy initialization of a class member using `std::call_once()` vs `std::mutex` (s2to7)
 - Example of actual usage of `std::call_once` (s2to8)

Protecting shared data with mutexes

- General guidelines for avoiding deadlocks
 - Avoid nested locks
 - Avoid calling user-supplied code while holding a lock
 - Acquire locks in a fixed order
- Extend these guidelines beyond locks

Protecting shared data with mutexes

- Structuring code for protecting shared data
 - Not as easy as slapping `std::lock_guard` everywhere
 - Spotting race conditions inherent in interfaces (Example: stack)
 - Option 1: Pass in a reference
 - Option 2: Require a no-throw copy constructor or move constructor
 - Option 3: Return a pointer to the popped item
 - Option 4: Provide both option 1 and either 2 or 3
- Example of definition of thread safe stack (**s2t13**)

Protecting shared data with mutexes

- Summary
 - How to use `std::mutex` and `std::lock_guard`
 - Avoiding deadlock with `std::lock`
 - Alternative data protection facilities like locking hierarchy or `std::call_once`
 - Example of a broken-by-design thread-(un)safe interface
- Next meeting:
 - Waiting for events
 - `std::future`

The end

- Any questions?